

ApexPilot AI — User Guide & Visual Walkthrough

1. Tailoring Answers with Candidate Resume & Job Description (JD)

1-Click File Loading (.pdf, .txt, .md)

⚡ Local Routing

🌐 Online API Pool

🔧 Custom Instructions

📄 Resume Context

🔒 Stealth Storage

Candidate Resume & Experience Context:

📁 Load Resume (.pdf, .txt, .md)

Senior Distributed Systems Engineer (8+ years experience).

- Architected ultra-low latency event streaming backends using Python, Go, Kafka, Redis, and C++.
- Led migration of core monolith to Kubernetes, cutting compute costs by 40% (\$320k/yr).
- Scaled real-time order matching pipeline to 5M events/sec with sub-20ms p99 latency.
- Expert in distributed consensus (Raft/Paxos), database internals (PostgreSQL, Cassandra), and eBPF.

Target Job Description (JD):

📁 Load JD (.pdf, .txt, .md)

Staff Software Engineer - Infrastructure & Distributed Systems.


Requirements: Deep knowledge of high-throughput messaging (Kafka/Pulsar), microservices latency optimization, and concurrency. Must possess strong STAR behavioral leadership experience and architectural design skills.

💡 ApexPilot AI tailors STAR behavioral stories, system designs, and talking points specifically to your authentic resume projects and the compan

Cancel

Save Apply

How Dual-Context Ingestion Drives Responses



2. Master Stealth HUD & Coding Copilot

SHIELD: EXCLUDED FROM ZOOM / MEET / TEAMS

Teleprompter

Click-Thru: OFF

StealthCoder

STAR Behavioral

System Design

HuddleMate

Snip LeetCode Screen (Ctrl+Alt+S)

Copy Clean Code

Saved PDFs

Opacity:

⚡ Intuition & Approach

- Use a **Hash Map** combined with a **Doubly Linked List** to achieve **O(1)** time complexity for both get and put.
- The hash map stores keys mapped directly to node references in the linked list. The doubly linked list maintains the temporal usage order (Most Recently Used at head, Least Recently Used at tail).
- Time Complexity:** **O(1)** amortized for all operations.
- Space Complexity:** **O(Capacity)** auxiliary memory.

📄 Optimal Production Code (Python)

```

class Node:
    __slots__ = ('key', 'val', 'prev', 'next')
    def __init__(self, key=0, val=0):
        self.key, self.val = key, val
        self.prev = self.next = None

class LRUCache:
    def __init__(self, capacity: int):
        self.capacity = capacity
        self.cache = {} # key -> Node
        self.head, self.tail = Node(), Node()
        self.head.next, self.tail.prev = self.tail, self.head

    def _remove(self, node: Node):
        node.prev.next = node.next

```

Design a high-throughput LRU Cache with O(1) get & put operations.

⚡ Solve

Ctrl+Alt+S Snip

Ctrl+Alt+H Panic

Ctrl+Alt+T Teleprompter

Ctrl+Alt+C Copy

Drag borders to resize

3. Local LLMs & Speculative Race Routing

⚡ Local Routing

🌐 Online API Pool

🔥 Custom Instructions

📄 Resume Context

🔒 Stealth Storage

Primary / Local LLM Server

Primary Provider:

ollama

Local Base URL:

http://127.0.0.1:11434

Local Model:

qwen2.5-coder:7b

⚡ Test Primary Connection

✅ Ollama Online (4 models available)

Fastest-First Race 429 Failover

☒ ⚡ Fastest-First Race Mode (Speculative Duel: Local vs Online LLM parallel race)

☒ 🔄 Automatic 429 Rate-Limit Failover (Rotates to next API key/provider if rate-limited)

StealthCoder Language:

Go

Cancel

Save Apply

4. Unlimited Online LLM API Pool

⚡ Local Routing

🌐 Online API Pool

🔥 Custom Instructions

📄 Resume Context

🛡️ Stealth Storage

Configure Online LLM API Keys for Auto-Failover & Fastest-First Duel:

OpenAI (ChatGPT)

☒ Enable API Key: Model: gpt-4o

xAI (Grok)

☒ Enable API Key: Model: grok-2

Groq (Ultra-Fast LPU 300+ tok/s)

☒ Enable API Key: Model: llama-3.3-70b-versatile

Anthropic (Claude)

☐ Enable API Key: sk-ant-... Model: gpt-3.5-sonnet-20241022

DeepSeek

☐ Enable API Key: sk-... Model: deepseek-chat

Google Gemini

☐ Enable API Key: AlzaSy... Model: gemini-2.0-flash

Cancel

Save Apply

5. Webcam Teleprompter & LeetCode Screen Snip

🛡️ STEALTH

◀ 1/3 ▶

Scaled Apache Kafka cluster to 5M events/sec with zero packet loss during peak market volatility.

Main HUD

 [SCREEN SNIP REGION DETECTED — NATIVE WINDOWS MEDIA OCR (<25ms)]

146. LRU Cache (Medium)

Design a data structure that follows the constraints of a Least Recently Used (LRU) cache.

Implement the `LRUCache` class:

- `LRUCache(int capacity)` Initialize the LRU cache with positive size capacity.
- `int get(int key)` Return the value of the key if key exists, otherwise return -1.
- `void put(int key, int value)` Update or insert key-value pair. Evict LRU if capacity exceed

 Press **Ctrl + Alt + S** anywhere. Drag red crosshair over question. Text extracts in

6. Q&A Instant Cache & Local PDF Archiving

ApexPilot AI - Saved Interview Archive & Vector PDF

Interview_Report_LRU_Cache_20260923_120035.pdf

• Mode: StealthCoder • Model: Groq Llama-3.3-70B • Latency: 0ms (Cache Hit) • Status: Archived

Question: Design a high-throughput Least Recently Used (LRU) Cache in Python.

Candidate Profile Match: Aligned with Senior Distributed Systems resume context.
Time Complexity: O(1) • Space Complexity: O(Capacity)
Full Code & Speakable Explanation Archived Locally.

Location: d:\interAI\saved_interviews\ • Click ' Saved PDFs' in HUD to open anytime.

7. Global Hotkeys & Controls Reference

Key	Action	Notes
Alt+Q	Quick Search	Searches through the current document for the specified text.
Alt+W	Window Management	Opens a panel to manage open windows and tabs.
Alt+E	Export Document	Exports the current document to a PDF or other format.
Alt+P	Print Document	Prints the current document.
Alt+S	Save Document	Saves the current document.
Alt+R	Refresh Content	Refreshes the content of the current view.
Alt+V	View Settings	Opens a panel to adjust view settings like theme and font.

8. Getting Started